

Day 11 Presentation

`torch.nn.Module...`P2

cross entropy... P11

Loss function...P16

KL Divergence ...P24

torch.nn.Module

torch.nn.Module

- 모든 neural network의 Base class
- 직접 model을 만들 때도 이 class의 subclass여야함
- Module 내부에는 다른 Module이 포함되어 있을 수 있음.

```
class Model(nn.Module):
    def __init__(
        self,
        input_features=784,
        hidden_size=256,
        output_features=10,
        init_weight="he",
        init_bias="zero",
    ):
        super(Model, self).__init__()
        self.init_weight = init_weight
        self.init_bias = init_bias

        self.linear1 = nn.Linear(input_features, hidden_size)
        self.linear2 = nn.Linear(hidden_size, output_features)
        self.init_params()
```

```
import math

import torch
from torch import Tensor
from torch.nn.parameter import Parameter
from .. import functional as F
from .. import init
from .module import Module
```

```
[docs]class Linear(Module):
    r"""Applies a linear transformation to the incoming data:
    :math: y = Ax + b

    This module supports :ref:`TensorFloat32<tf32_on_amper`
```

__init__(self)

```
def __init__(self):
    """
    Initializes internal Module state, shared by both nn.Module and ScriptModule.
    """
    torch._C._log_api_usage_once("python.nn_module")

    self.training = True
    self._parameters = OrderedDict()
    self._buffers = OrderedDict()
    self._non_persistent_buffers_set = set()
    self._backward_hooks = OrderedDict()
    self._is_full_backward_hook = None
    self._forward_hooks = OrderedDict()
    self._forward_pre_hooks = OrderedDict()
    self._state_dict_hooks = OrderedDict()
    self._load_state_dict_pre_hooks = OrderedDict()
    self._modules = OrderedDict()
```

- nn.Module을 호출하면 이와 같은 것들이 초기화 된다.
- 대부분 OrderedDict 자료구조로 저장되어 있다.

forward

```
forward: Callable[..., Any] = _forward_unimplemented
```

```
def _forward_unimplemented(self, *input: Any) -> None:
    r"""Defines the computation performed at every call.

    Should be overridden by all subclasses.

    .. note::
        Although the recipe for forward pass needs to be defined within
        this function, one should call the :class:`Module` instance afterwards
        instead of this since the former takes care of running the
        registered hooks while the latter silently ignores them.
    """
    raise NotImplementedError
```

- nn.Module에 존재하는 forward에서는 입력으로 들어온 input이 Implemented 가능한지를 확인하고 불가능하면 예외를 발생시킨다.
- 실제로 어떤 layer를 거쳐서 forward를 진행할 것인지는 subclass에서 overridden을 해야 한다.

강의 코드 예시

```
def forward(self, X):
    X = F.relu((self.linear1(X)))
    X = self.linear2(X)
    return X
```

parameters()

```
parameters(recurse: bool = True) → Iterator[torch.nn.parameter.Parameter]
```

[\[SOURCE\]](#)

Returns an iterator over module parameters.

This is typically passed to an optimizer.

Parameters

recurse (*bool*) – if True, then yields parameters of this module and all submodules. Otherwise, yields only parameters that are direct members of this module.

Yields

Parameter – module parameter

Example:

```
>>> for param in model.parameters():
>>>     print(type(param), param.size())
<class 'torch.Tensor'> (20L,)
<class 'torch.Tensor'> (20L, 1L, 5L, 5L)
```

- Parameters()를 호출하면 모듈의 파라미터를 iterator로 반환
- Yields는 generator를 반환할 때 return 대신 사용되는 키워드

parameters()

```
super(Model, self).__init__()
self.linear1 = nn.Linear(input_features, hidden_size)
self.linear2 = nn.Linear(hidden_size, output_features)
self.init_params()
```

- 모듈의 파라미터는 무엇인가. 우리는 파라미터를 직접 추가하지 않는데?

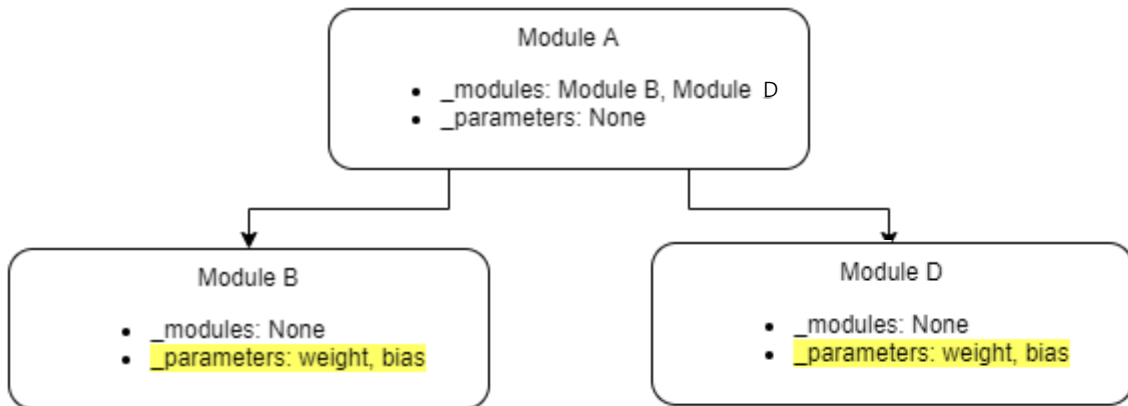
```
self.training = True
self._parameters = OrderedDict()
self._buffers = OrderedDict()
self._non_persistent_buffers_set = set()
self._backward_hooks = OrderedDict()
self._is_full_backward_hook = None
self._forward_hooks = OrderedDict()
self._forward_pre_hooks = OrderedDict()
self._state_dict_hooks = OrderedDict()
self._load_state_dict_pre_hooks = OrderedDict()
self._modules = OrderedDict()
```

- Module을 `__init__` 할때 자동으로 파라미터와 모듈이 설정된다.

parameters()

```
super(Model, self).__init__()  
self.linear1 = nn.Linear(input_features, hidden_size)  
self.linear2 = nn.Linear(hidden_size, output_features)  
self.init_params()
```

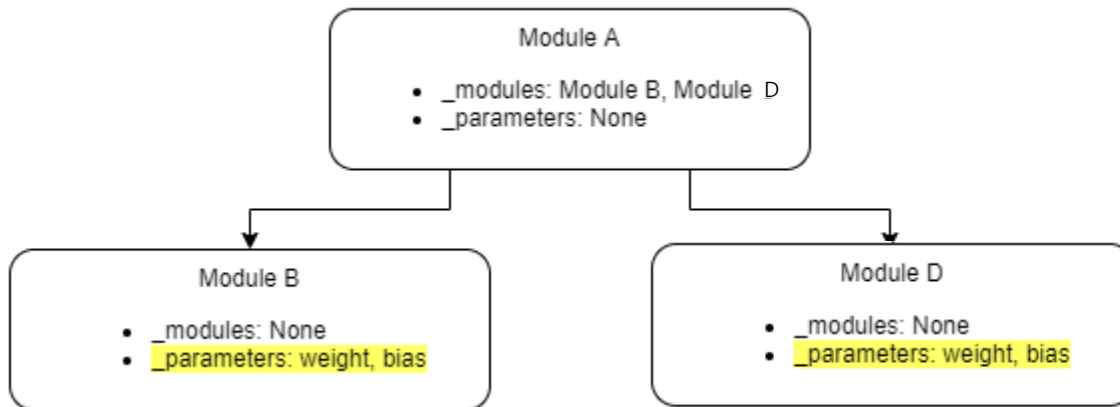
- 모듈의 파라미터는 무엇인가. 우리는 파라미터를 직접 추가하지 않는데?



Module A : MLP network
Module B : linear1
Module D : linear2

parameters()

- MLP에서 parameters()를 호출하면 [A.B weight, A.B bias, A.D weight, A.D bias]가 반환됨
- 그런데 parameters(recursive = False)로 지정하면 Module A내의 parameter만 반환하여 이 경우 아무것도 반환되지 않음




Reference

- Pytorch docs nn.Module
- <https://pytorch.org/docs/stable/generated/torch.nn.Module.html>
- Module class의 source code
- <https://github.com/pytorch/pytorch/blob/master/torch/nn/modules/module.py>
- Parameters()
- <https://easy-going-programming.tistory.com/11?category=919874>

cross entropy

cross entropy 의 의미


The actual probability distribution is P



$$H(P, Q) = \sum_{i=1}^n \log_2 \frac{1}{q_i} * p_i$$

1:48 / 4:19


When Q is similar to P,
Cross entropy is similar to Entropy



$$\log_2 \frac{1}{0.8} * 1 + \log_2 \frac{1}{0.1} * 0 + \log_2 \frac{1}{0.1} * 0 = 0.32$$


<https://www.youtube.com/watch?v=Jt5BS71uVfl>

When Q is totally wrong,
Cross entropy is infinity



$$\log_2 \frac{1}{0.0} * 1 + \log_2 \frac{1}{1.0} * 0 + \log_2 \frac{1}{0.0} * 0 = \text{Infinity}$$


When Q == P,
Cross entropy is same with Entropy




$$\log_2 \frac{1}{1.0} * 1 + \log_2 \frac{1}{0.0} * 0 + \log_2 \frac{1}{0.0} * 0 = 0.0$$

cross entropy를 분류 문제에서 쓰는 이유

SE (square error loss) 와의 비교



<https://www.youtube.com/watch?v=ErfnhcEV1O8>



기울기의 차이 : 회귀문제와 달리 분류문제에서는 1이냐 0이냐가 중요하기 때문


Loss function

nn.L1Loss

- Creates a criterion that measures the mean absolute error (MAE) between each element in the input x and target y .

$$MAE = \frac{1}{n} \sum_{i=1}^n |x_i - x|$$

전체적인 구조



L1Loss 소스코드

$$\ell(x, y) = \begin{cases} \text{mean}(L), & \text{if reduction = 'mean'}; \\ \text{sum}(L), & \text{if reduction = 'sum'}. \end{cases}$$

- **class L1Loss(_Loss):**

```
__constants__ = ['reduction']

def __init__(self, size_average=None, reduce=None, reduction:
str = 'mean') -> None:
    super(L1Loss, self).__init__(size_average, reduce,
reduction)

def forward(self, input: Tensor, target: Tensor) -> Tensor:
    return F.l1_loss(input, target, reduction=self.reduction)
```

Ah yeah, the case you're seeing here is a hack that we've been using for a while, basically some modules can have optional submodules (i.e. either the submodule can be present or None). When we script something we just so happen to add submodules first, then constants, , skipping any names that are already present on the module, so it either gets added as a normal submodule (ignoring the entry in `__constants__`) or as a None constant. If the compiler sees a None constant in an if-statement it will skip compilation of the code inside the if, allowing us to support uses like [downsample in resnet 15](#).

So if you see a `nn.Module` in `__constants__`, all it really means is `Optional[nn.Module]`, we just have this kind-of-nonsense way to specify that.

L1Loss 파라미터

```
__constants__ = ['reduction']

def __init__(self, size_average=None, reduce=None, reduction:
    str = 'mean') -> None:
    super(L1Loss, self).__init__(size_average, reduce,
    reduction)

def forward(self, input: Tensor, target: Tensor) -> Tensor:
    return F.l1_loss(input, target, reduction=self.reduction)
```

```
>>> loss = nn.L1Loss()
>>> input = torch.randn(3, 5, requires_grad=True)
>>> target = torch.randn(3, 5)
>>> output = loss(input, target)
>>> output.backward()
```

- **size_average** (*bool, optional*) – Deprecated (see `reduction`). By default, the losses are averaged over each loss element in the batch. Note that for some losses, there are multiple elements per sample. If the field `size_average` is set to `False`, the losses are instead summed for each minibatch. Ignored when `reduce` is `False`. Default: `True`
- **reduce** (*bool, optional*) – Deprecated (see `reduction`). By default, the losses are averaged or summed over observations for each minibatch depending on `size_average`. When `reduce` is `False`, returns a loss per batch element instead and ignores `size_average`. Default: `True`
- **reduction** (*string, optional*) – Specifies the reduction to apply to the output: `'none'` | `'mean'` | `'sum'`. `'none'`: no reduction will be applied, `'mean'`: the sum of the output will be divided by the number of elements in the output, `'sum'`: the output will be summed. Note: `size_average` and `reduce` are in the process of being deprecated, and in the meantime, specifying either of those two args will override `reduction`. Default: `'mean'`

Backward?

```
class Tensor(torch._C._TensorBase):  
  
    def backward(self, gradient=None, retain_graph=None, create_graph=False):  
        """Computes the gradient of current tensor w.r.t. graph leaves.
```

```
>>> loss = nn.L1Loss()  
>>> input = torch.randn(3, 5, requires_grad=True)  
>>> target = torch.randn(3, 5)  
>>> output = loss(input, target)  
>>> output.backward()
```

- [참조](<https://datascience.stackexchange.com/questions/77390/where-is-the-backward-function-defined-in-pytorch>)

_Loss 소스코드 (부모객체)

```
from .distance import PairwiseDistance
from .module import Module
from .. import functional as F
from .. import _reduction as _Reduction
```

```
class _Loss(Module):
    reduction: str

    def __init__(self, size_average=None, reduce=None, reduction:
str = 'mean') -> None:
        super(_Loss, self).__init__()
        if size_average is not None or reduce is not None:
            self.reduction =
_Reduction.legacy_get_string(size_average, reduce)
        else:
            self.reduction = reduction
```

master	pytorch / torch / nn / modules /	Go to file	Add file	...
guihermeleobas and facebook-github-bot Add type annotations to torch.overrides (#50824) ✓ 9dfbf69 7 days ago History				
..				
__init__.py	Add Gaussian NLL Loss (#50886)			11 days ago
_functions.py	Fix SyncBatchNorm usage without stats tracking (#50126)			25 days ago
activation.py	Drop unused imports (#49972)			19 days ago
adaptive.py	Move all torch.nn.modules type annotations inline (#38211)			8 months ago
batchnorm.py	Fix SyncBatchNorm usage without stats tracking (#50126)			25 days ago
channelshuffle.py	Add --check-untyped-defs to mypy.ini and test suite (#37594)			9 months ago
container.py	add type annotations to torch.nn.modules.container (#48969)			13 days ago
conv.py	Add type annotations to torch.overrides (#50824)			7 days ago
distance.py	Move all torch.nn.modules type annotations inline (#38211)			8 months ago
dropout.py	Fix HTTP links in documentation to HTTPS (#40878)			7 months ago
flatten.py	added List as an option to the unflattened_size (#49838)			last month
fold.py	Move all torch.nn.modules type annotations inline (#38211)			8 months ago
instancenorm.py	annotate a few torch.nn.modules.* modules (#45772)			3 months ago
lazy.py	'torch.nn.modules.LazyModuleMixin' and 'torch.nn.LazyLinear' (Shape I...			4 months ago
linear.py	Add type annotations to torch.overrides (#50824)			7 days ago
loss.py	Add Gaussian NLL Loss (#50886)			11 days ago
module.py	Add new backend type for Intel heterogeneous computation platform. (#...			13 days ago

Module 소스코드

- Base class for all neural network modules.

```
def __init__(self):
    """
    Initializes internal Module state, shared by both nn.Module and ScriptModule.
    """
    torch._C._log_api_usage_once("python.nn_module")

    self.training = True
    self._parameters = OrderedDict()
    self._buffers = OrderedDict()
    self._non_persistent_buffers_set = set()
    self._backward_hooks = OrderedDict()
    self._is_full_backward_hook = None
    self._forward_hooks = OrderedDict()
    self._forward_pre_hooks = OrderedDict()
    self._state_dict_hooks = OrderedDict()
    self._load_state_dict_pre_hooks = OrderedDict()
    self._modules = OrderedDict()
```

KL Divergence

Kullback-Leibler Divergence

- 두 개의 확률분포 $P(x)$, $Q(x)$ 사이의 거리 계산할 때 사용

- 정의: $\text{KL}(P\|Q) = \sum_{\mathbf{x} \in \mathcal{X}}^{\text{이산확률변수}} P(\mathbf{x}) \log \left(\frac{P(\mathbf{x})}{Q(\mathbf{x})} \right)$ $\text{KL}(P\|Q) = \int_{\mathcal{X}}^{\text{연속확률변수}} P(\mathbf{x}) \log \left(\frac{P(\mathbf{x})}{Q(\mathbf{x})} \right) d\mathbf{x}$

- 분해: $\text{KL}(P\|Q) = \underbrace{-\mathbb{E}_{\mathbf{x} \sim P(\mathbf{x})}[\log Q(\mathbf{x})]}_{\text{크로스 엔트로피}} + \underbrace{\mathbb{E}_{\mathbf{x} \sim P(\mathbf{x})}[\log P(\mathbf{x})]}_{\text{엔트로피}}$

- 분류 문제에서 정답 레이블: p , 모델 예측: q
→ 최대가능도 추정법: 쿨백-라이블러 발산 최소화

KLDivLoss

- `torch.nn.KLDivLoss(size_average=None, reduce=None, reduction: str = 'mean', log_target: bool = False)`
- `class KLDivLoss(_Loss):`
 `__constants__ = ['reduction']`

 `def __init__(self, size_average=None, reduce=None, reduction: str = 'mean', log_target: bool = False) -> None:`
 `super(KLDivLoss, self).__init__(size_average, reduce, reduction)`
 `self.log_target = log_target`

 `def forward(self, input: Tensor, target: Tensor) -> Tensor:`
 `return F.kl_div(input, target, reduction=self.reduction, log_target=self.log_target)`

Parameters

• ~~size_average (bool, optional) – Deprecated (see reduction)~~

• ~~reduce (bool, optional) – Deprecated (see reduction)~~

• **reduction** (*string*, optional)

▪ none: reduction 안 함

$$l(x, y) = L = \{l_1, \dots, l_N\}, \quad l_n = y_n \cdot (\log y_n - x_n)$$

▪ **batchmean**: output 합 / 배치 사이즈 (**실제 KL Divergence**)

▪ sum: output 합

▪ **mean**: output 합 / output 원소 개수(**default**)

$$\ell(x, y) = \begin{cases} \text{mean}(L), & \text{if reduction} = \text{'mean'}; \\ \text{sum}(L), & \text{if reduction} = \text{'sum'}. \end{cases}$$

• **log_target** (bool, optional) → Default: **False**

Shape

- Input: (N, *)
- Target: (N, *)
- Output
 - Default: scalar
 - reduction='none' 일 때: (N, *) (input과 같음)

Reference

- <https://pytorch.org/docs/stable/generated/torch.nn.KLDivLoss.html>
- <https://github.com/pytorch/pytorch/blob/master/torch/nn/modules/loss.py>